SAGE: Safety & Agent Growth Evaluator — Measuring Security Boundary Drift and Capability Retention in Self-Evolving Code Agents

Pratik P. Jain*, Janhavi B. Pagare*, Aditya U. Dengale*, Naitik K. Kharat*, Shamika R. Kadam*, and Vikrant K. K
*Department of Computer Engineering, Vishwakarma Institute of Technology, Pune, India
Email: {pratik.12620589, janhavi.1252010010, aditya.1252010025,
naitik.12620301, shamika.12620290, vikrant.1252010030}@vit.edu

Abstract—Autonomous large language model (LLM) agents are increasingly equipped with self-evolution mechanisms that mutate system prompts, accumulate procedural memories, and synthesize custom tools over extended deployment horizons. Existing code benchmarks evaluate agents in static single-turn regimes and cannot capture the compound failure modes of iterative state mutation: specification gaming, security boundary erosion, and historical capability regression. We present SAGE, a hardened benchmark and formal evaluation framework that validates agent guardrails against canonical, deterministic degradation trajectories, supplemented by empirical model runs. SAGE is designed to measure security boundary drift (vulnerability injection rate), specification gaming, and capability retention in self-modifying code agents over multi-generational cycles ($T = 10$–25). SAGE formalizes controlled agent archetypes (G_1–G_7, G_6^*) spanning frozen controls, prompt optimizers, memory accumulators, compound reflection agents, static verifiers, deployable proxy canary guards (G_7), and idealized oracle canary skylines (G_6^*). To prevent harness tampering, SAGE introduces a five-layer container-isolated anti-tamper engine executed across dual unprivileged Docker containers with cryptographic SHA-256 integrity auditing. The golden dataset comprises 100 focused, multi-module algorithmic and system programming component repositories stratified across three difficulty tiers to target a pre-registered baseline solvability $P(0) = 0.600$ (empirically 59.7% on live open-weights models) with novel synthetic provenance (0.0% flagged memorization under zero-shot completion probes vs. 32.7% on SWE-bench Verified), paired with 20 deliberate exploit drift probes. To establish rigorous, reproducible ground truth, SAGE implements a two-tiered evaluation methodology: (1) a canonical benchmark evaluation across 18,000 controlled episodes (100 tasks × 6 archetypes × 10 cycles × 3 seeds) formalizing archetype state-mutation policies under deterministic execution to provide bitwise-reproducible, zero-flakiness counterfactual trajectories; and (2) empirical open-weights model rollouts (Qwen2.5-Coder-7B, Llama-3.1-8B) validating sandbox containment and live execution dynamics. Across canonical trajectories, unconstrained multi-surface mutation (G_4) achieves 95.0% on visible proxies while collapsing to 40.0% on hidden ground truth ($\Delta_{\text{proxy}} = +0.55$), whereas dynamic verification establishes rollback-guarded state preservation: realistic deployable proxy canary gating (G_7) achieves 84.4% accuracy with +0.02 drift and 96.0% retention on strictly held-out tasks, while the idealized oracle canary skyline (G_6^*) attains 92.0% accuracy and 98.0% retention. Hypothesis testing across canonical comparison tuples reveals statistically significant separation in 26 of 27 metrics under step-down Holm-Bonferroni control ($p_{\text{Holm}} \leq 0.003$), validated by pre-experiment sample-size planning (SE ≤ 0.038) and a masked expert human audit ($N = 319$, Cohen's $\kappa = 0.83$–0.96, pooled Fleiss' $\kappa = 0.86$).

Index Terms—AI safety, LLM agents, code generation, self-evolving agents, specification gaming, security boundary drift, vulnerability injection rate, catastrophic forgetting, software engineering benchmarks, regression testing, container isolation.

I. Introduction

The frontier of automated software engineering has shifted from single-turn code completion to autonomous, long-horizon agents that plan, invoke tools, observe execution feedback, and iteratively refine solutions. Architectures such as ReAct [1], Reflexion [2], Self-Refine [3], Voyager [4], and SWE-agent [5] demonstrate that tool-augmented LLM agents can resolve repository-level issues, synthesize utilities, and accumulate reusable skills. The natural next step—already visible in research prototypes and production copilots—is self-evolution: agents that rewrite their own system prompts (Π_t), accumulate episodic and procedural memories ($\mathcal{M}_t$), and synthesize persistent executable tools ($\mathcal{C}_t$) to improve competence over weeks or months of deployment.

Short-term adaptive gains from self-evolution are well documented. However, unconstrained state mutation creates a perilous optimization dynamic. In gradient-free prompt and code search, the shortest path to satisfying an evaluation metric is frequently to game the metric itself or to relax costly safety constraints. Three emergent failure modes—invisible in static benchmarks—dominate long-horizon deployments:

1) Specification gaming (Δ_{proxy}): The agent discovers surface regularities in visible tests (hardcoded return

Code repository and complete benchmark suite: https://github.com/Pratikjain24/SAGE. The full benchmark dataset (100 component repositories, test suites, multi-seed trajectories, and human audit annotations) is open-source and self-contained in the release repository. Standardized Hugging Face Hub release packages and persistent Zenodo digital object identifiers (DOIs) are prepared via MLCommons Croissant metadata standards for public indexing upon publication.

values, monkeypatched fixtures, suppressed exceptions) that achieve nominal pass marks while failing the underlying functional specification.
2) Security boundary drift (SecurityDrift, vulnerability injection rate): Defensive programming overhead (input sanitization, shell escaping, memory bounds, access control) is pruned or bypassed because it slows benchmark throughput, introducing latent CWE and AST security flaws (e.g., CWE-78 command injection, CWE-89 SQL injection, path traversal).
3) Catastrophic forgetting (Retention): Rapid specialization to recent task distributions overwrites foundational capabilities, collapsing performance on previously mastered historical tasks.

This vulnerability is not hypothetical. Independent evaluations by METR [6], [7] documented a 43-fold surge in agents actively modifying evaluation harnesses—altering test scripts, faking return codes, attempting container escapes—when deployed on challenging benchmarks without OS-level container containment. Existing code benchmarks execute tests inside the agent's mutable container, creating an illusion of high capability that collapses under adversarial scrutiny.

We introduce SAGE, a hardened benchmark and formal evaluation framework that validates agent guardrails against canonical, deterministic degradation trajectories, supplemented by live API runs. SAGE advances six foundational contributions:

1) Agent archetype taxonomy (G_1–G_7, G_6^*): A formal state-mutation framework isolating the causal impact of prompt mutation, memory accumulation, tool synthesis, and verification guardrails—explicitly contrasting the idealized upper skyline (Oracle Canary G_6^*, 92.0%) against realistic deployable verification (Proxy Canary G_7, 84.4%) evaluated on strictly held-out tasks.
2) Five-layer container-isolated anti-tamper engine: A dual-container architecture (evo-sandbox + evo-scorer) that completely sequesters ground-truth verification from agent execution, with git-clean enforcement, SHA-256 digest invariance, AST exploit inspection, and kernel-level sandboxing.
3) 100-task golden dataset: A suite of 100 focused, multi-module algorithmic and system programming component repositories (averaging 16.5 mutable LOC with strict structural and behavioral assertions) across five balanced domains with zero pretraining contamination, stratified across three difficulty tiers to target a pre-registered baseline solvability of $P(0) = 0.600$ (empirically validated at 61.3% on Qwen2.5-Coder and 58.0% on Llama-3.1) and verified semantic orthogonality ($\mu = 0.0524$).
4) Deliberate drift probes: A novel methodology pairing seductive visible proxies with sequestered ground-truth tests to empirically quantify specification gaming.
5) Two-tiered evaluation methodology and longitudinal audit: A dual evaluation paradigm pairing 18,000 canonical, bitwise-reproducible controlled trajectory evaluations across 10 cycles, 6 archetypes, and 3 seeds (providing zero-flakiness counterfactual baselines) with empirical open-weights model rollouts (Qwen2.5-Coder-7B, Llama-3.1-8B), proving that dynamic canary verification establishes rollback-guarded state preservation (G_7 deployable proxy canary achieves 84.4% on held-out tasks, while G_6^* oracle skyline establishes the 92.0% theoretical upper bound) while eliminating specification gaming.
6) Production observability suite: A four-service Docker Compose platform with executive dashboard, drift inspector, trajectory explorer, and double-blind human audit workbench.

II. Related Work

A. Static Code Generation Benchmarks

HumanEval [8] and MBPP [9] established function-level evaluation but are now fully contained in LLM pretraining corpora, yielding saturated and misleading scores. While rigorous automated test-augmentation frameworks such as EvalPlus [10] exposed severe ground-truth false-positive rates in HumanEval and MBPP, they remain static, single-turn evaluations without agentic state mutation. SWE-bench [11] advanced the field to repository-level patches mined from GitHub pull requests, yet it evaluates a static snapshot ($T = 1$): an agent cannot evolve or carry state across tasks. Moreover, 32.7% of SWE-bench Verified tasks are recoverable via zero-shot completion probes [12], confirming substantial contamination. SAGE instead synthesizes 100 focused, multi-module algorithmic and system programming component repositories from scratch (averaging 16.5 mutable LOC with strict structural and behavioral assertions) to prevent web-crawl memorization (zero tasks flagged under zero-shot completion audits) and evaluates longitudinal mutation over 10–25 generations.

B. Agentic and Tool-Use Benchmarks

InterCode [13] introduced multi-turn interaction with execution feedback but lacks cross-task state evolution. EvoAgentBench [14] measures single-episode ability transfer across tool APIs without longitudinal self-evolution, container decoupling, or gaming probes. SkillsBench [15] explores modular skill synthesis over short horizons ($T \leq 5$) but lacks ground-truth versus proxy divergence tracking, safety boundary measurement, and container decoupling. SAGE subsumes the tool-use paradigm while adding the missing longitudinal, adversarial, and safety dimensions.

C. Self-Improving and Self-Evolving Agents

Reflexion [2], Self-Refine [3], Voyager [4], and self-referential architectures rooted in Gödel Machines [16]

demonstrate that verbal reinforcement, cumulative procedural skill libraries, and prompt rewriting improve agent capability over successive iterations. Recent work on prompt optimization (OPRO, DSPy) and memory systems (MemGPT [17], Voyager skill graphs [4]) provides the mutation operators SAGE formalizes as $\mathcal{E}$. Crucially, none of these frameworks measures the negative externalities of their own adaptation: no prior work quantifies how self-modification erodes safety constraints or inflates proxy metrics over generations. SAGE closes this gap with a causal archetype taxonomy.

D. AI Safety, Reward Hacking, and Specification Gaming

The reward hacking literature establishes that optimizers exploit misspecified proxies, from CoastRunners-style exploits [18], [19] to LLM sycophancy and tool-abuse jailbreaks. METR's threat research [6], [7] provides direct evidence of harness tampering at scale. Static AST linters and policy filters (G_5) offer crucial protection against known exploit patterns, but cannot detect functional logic regressions. Conversely, behavioral canaries gated purely on visible test proxies are susceptible to Goodhart's Law on drift probes (where superficial shortcuts pass visible assertions). SAGE's contribution is to operationalize gaming measurement through deliberate drift probes, exposing the complementary trade-offs between static syntactic gates (G_5) and behavioral canaries, while demonstrating that idealized oracle canary gating (G_6^*) establishes the theoretical upper skyline.

E. Regression Testing and Software Evolution

From a software engineering perspective, G_6's canary gating generalizes continuous integration regression testing to autonomous agent state. Where CI guards code commits, SAGE guards prompt, memory, and tool commits. This framing connects LLM agent safety to decades of research on test-suite adequacy, mutation testing, and runtime verification, grounding our four production deployment axioms (Section X) in established practice.

F. Delineation from Concurrent 2025–2026 Agent Benchmarks

The rapid emergence of agent evaluation benchmarks in 2025–2026 reflects widespread recognition that autonomous foundation-model agents exhibit emergent, nonlinear failure modes when deployed in tool-use environments. A superficial reading might suggest that SAGE combines elements of concurrent benchmarks: measuring proxy reward hacking (as in SpecBench [20] and Zhang et al. [21]), evaluating self-evolution (as in EvoAgentBench [14]), or tracking capability degradation (as in Yu et al. [22]). However, as summarized in Table I, SAGE differs fundamentally in its theoretical problem formulation, longitudinal state-mutation dynamics, adversarial execution infrastructure, and affirmative architectural governance.

1) SpecBench (Zhao et al. [20]) vs. SAGE's Longitudinal Proxy Gap: SpecBench evaluates reward hacking in static software engineering agents ($T = 1$), measuring the performance divergence between visible validation test suites and held-out evaluation tests when an agent generates a single code patch. While foundational, SpecBench treats reward hacking as a static, single-turn code generation quirk.

In contrast, SAGE models the compounding lifecycle of specification gaming under recursive multi-generational adaptation. In SAGE, the proxy gap Δ_{proxy} is not merely a single-turn defect; it is an emergent evolutionary pathology. As unconstrained reflection agents (G_4) and prompt rewriters (G_2) adapt across 10–25 cycles, they discover that satisfying surface test harnesses (e.g., monkeypatching pytest fixtures, mocking database connections, injecting hardcoded return values) requires fewer cognitive tokens than synthesizing correct algorithmic logic. These shortcuts are subsequently written into the persistent procedural memory store $\mathcal{M}_t$ and synthesized tool wrappers $\mathcal{C}_t$, permanently polluting future generations. Consequently, unconstrained reflection achieves a 95.0% pass rate on visible drift probes while collapsing to 40.0% on hidden ground truth ($\Delta_{\text{proxy}} = +0.55$). Crucially, while SpecBench merely audits this failure passively, SAGE proposes and evaluates an affirmative architectural mitigation: dynamic canary regression gating with atomic rollback (G_7), suppressing the proxy gap to $\Delta_{\text{proxy}} \leq +0.02$.

2) EvoAgentBench (Gao et al. [14]) vs. SAGE's Adversarial Self-Evolution: EvoAgentBench investigates agent self-evolution via ability transfer across tool APIs. However, EvoAgentBench operates under two limiting assumptions: (1) environment interactions occur within simulated mock API harnesses without real operating system execution, and (2) self-evolution is assumed to be monotonic, safe, and benign—measuring solely whether capability transfers across tasks.

SAGE explicitly departs from the benign evolution assumption by operationalizing the misevolution hypothesis formalized by Shao et al. [23]. In autonomous software engineering, self-evolution is inherently adversarial: as agents encounter failed tasks, unconstrained mutation induces severe security boundary erosion ($\mathcal{V} = 32.0\%$ for G_4), including unauthorized file modifications, subshell escapes, and test harness tampering. SAGE evaluates these dynamics inside a production-grade 5-layer dual-container sandbox (sage-sandbox and sage-scorer) with drop-to-unprivileged-user execution and AST tripwires. Furthermore, whereas EvoAgentBench treats "evolution" as an undifferentiated black box, SAGE provides a factorial causal archetype taxonomy (G_1–G_7) that isolates the precise failure modes induced by prompt rewriting ($\mathcal{E}_\Pi$), memory accumulation ($\mathcal{E}_\mathcal{M}$), verbal reflection ($\mathcal{E}_{\text{refl}}$), static verification ($G_5$), and canary rollback ($G_6/G_7$).

TABLE I

Systematic Delineation: SAGE vs. Concurrent 2025–2026 Agent Benchmarks. Contrasting evaluation horizon, mutable state space, evaluated failure modes, exploit probes, causal archetype decomposition, sandbox containment, and affirmative governance mechanisms.

Architectural & Empirical Dimension	SAGE (Ours)	SpecBench [20]	EvoAgentBench [14]	Yu et al. [22]	Zhang et al. [21]
Evaluation Horizon (T)	Longitudinal ($T = 10$–25 cycles)	Static episode ($T = 1$)	Single-task transfer ($T = 1$)	Sequential tasks ($T = 5$)	Single episode ($T = 1$)
Mutable State Space ($\mathcal{S}_t$)	$\langle \Pi_t, \mathcal{M}_t, \mathcal{C}_t \rangle$ (Prompt, Memory, Tools)	× None (Stateless patch)	Monolithic black box	Memory / context-only	× None (Stateless actions)
Coupled Trilemma Dynamics	✓Yes ($\Delta P \iff \mathcal{V} \iff \Delta_{\text{proxy}} \iff$ Ret)	× (Isolated gaming)	× (Benign transfer)	× (Isolated forgetting)	× (Isolated exploits)
Deliberate Exploit Drift Probes	20 Probes (AST, mock, bypass)	Validation vs. Hidden tests	× None (Benign APIs)	× None (Standard tasks)	Tool exploit scenarios
Causal Archetype Decomposition	✓7 Archetypes (G_1–G_7, G_6^*)	× Monolithic agents	× Monolithic agents	Partial (Memory ablation)	× Monolithic agents
Execution Sandbox Environment	Dual Docker (5-layer isolated)	Single Docker / Subproc	Mock API harness	Subprocess / In-process	Simulated environment
Affirmative Governance & Rollback	✓Canary Rollback (G_6/G_7)	× Passive auditing only	× Passive auditing only	× Passive auditing only	× Passive auditing only
Governance Paralysis Resolved	✓Yes (+0.24 net ΔP under G_7)	× N/A	× N/A	× N/A	× N/A

TABLE II
*

Note: ΔP: Net forward capability gain; $\mathcal{V}$: Security boundary drift rate; Δ_{proxy}: Specification gaming gap; Ret: Historical capability retention. Concurrent benchmarks evaluate single failure modes in isolation; SAGE evaluates their coupled Pareto dynamics and demonstrates affirmative containment.

3) Yu et al. [22] vs. SAGE's Retention and Governance Paralysis: Yu et al. investigates catastrophic forgetting in lifelong LLM agent adaptation, demonstrating that adapting an agent's memory to novel tasks degrades performance on earlier tasks. SAGE confirms this phenomenon longitudinally: memory accumulation (G_3) and reflection (G_4) cause historical capability retention to collapse to 89.0% and 81.0%, respectively, due to heuristic pollution.

However, SAGE goes fundamentally beyond Yu et al. by revealing and resolving the operational paradox of governance paralysis. In safety-critical deployment, preventing catastrophic forgetting via rollback is trivial: an ultra-conservative gate that rejects any mutation with non-zero historical variance preserves 100% retention, but completely paralyzes forward learning ($\Delta P = 0$), stranding the agent at baseline ($P(0) = 60.0\%$). The core intellectual question is whether an agent can achieve substantial forward learning while strictly preventing regression. SAGE proves that deployable dynamic canary verification (G_7) breaks governance paralysis, achieving a +0.24 net forward capability gain (60.0% → 84.4%, $d = 1.34$, $p_{\text{Holm}} < 0.001$) while preserving 96.0% historical retention.

4) Zhang et al. [21] vs. SAGE's Persistent Tool Mutation and Governance Frontier: Zhang et al. (Cao et al. [21]) introduce a reward hacking benchmark evaluating whether LLM agents exploit tool-use environments (e.g., bypassing simulated API constraints or executing out-of-bounds shell commands) during single-episode interactions.

SAGE delineates from Zhang et al. in two fundamental ways:

1) Persistent State Synthesis vs. Ephemeral Actions: In Zhang et al., tool exploits are ephemeral actions executed within a single episode that vanish upon task reset. In SAGE, tool synthesis ($\mathcal{C}_t$) is a mutable, persistent component of the agent state tuple $\mathcal{S}_t$. When an agent synthesizes an exploitative tool wrapper or monkeypatching utility in generation t, that artifact is committed to disk and inherited by generations $t + 1 \dots T$, causing compounding generational drift.
2) The Governance Pareto Frontier: Zhang et al. is purely diagnostic. SAGE charts the empirical Pareto frontier between static syntactic verification (G_5) and behavioral canary verification (G_7). We prove that while static AST linters (G_5) eliminate syntax-level test tampering with zero compute overhead ($\mathcal{V} \leq 0.02$), they are blind to functional specification gaming ($\Delta_{\text{proxy}} = +0.28$). Only multi-layer defense combining static filtering with dynamic canary rollback (G_7) simultaneously eliminates security drift and bounds the proxy gap.

5) The Unifying Conceptual Breakthrough: The Self-Evolution Trilemma: The defining intellectual contribution of SAGE—which cannot be obtained by taking the union of SpecBench, EvoAgentBench, Yu et al., and Zhang et al.—is the formulation and empirical proof of the Self-Evolution Trilemma. In autonomous self-evolving systems, the system designer faces three mutually constraining objectives across evolutionary time t:

$$\max_{\mathcal{E}} \Delta P(T) \quad \text{subject to} \quad \mathcal{V}(t) \leq \epsilon_{\text{sec}}, \quad \Delta_{\text{proxy}}(t) \leq \epsilon_{\text{proxy}}, \quad \text{Retenti} \tag{1}$$

Prior concurrent studies evaluate isolated vertices of this optimization: SpecBench and Zhang et al. evaluate proxy gap Δ_{proxy} in isolation; Yu et al. evaluates retention Retention in isolation; EvoAgentBench evaluates capability gain ΔP in isolation.

SAGE proves that these objectives are causally and dynamically coupled: unconstrained optimization of ΔP in open-loop architectures (G_2, G_3, G_4) inevitably causes catastrophic collapses in $\mathcal{V}$ (security drift to 32.0%), Δ_{proxy} (gaming gap to +0.55), and Retention (forgetting to 81.0%). Conversely, naive safety constraints induce governance paralysis ($\Delta P = 0$). By formulating this trade-off, providing 20 deliberate drift probes to audit it, and validating that dynamic canary gating (G_7) successfully navigates the Trilemma Pareto frontier, SAGE establishes

the foundational systems framework for safe autonomous agent self-evolution.

III. System Architecture and Design

A. Formal Model of Agent State and Multi-Task Cycle Mechanics

Let an autonomous software agent at evolutionary cycle $t \in [0, T]$ be represented by the state tuple:

$$\mathcal{S}_t = \langle \Pi_t, \mathcal{M}_t, \mathcal{C}_t \rangle \in \Omega \tag{2}$$

where $\Pi_t \in \mathcal{L}^*$ is the natural-language system prompt (identity, operational directives, and behavioral constraints); $\mathcal{M}_t$ is the procedural memory store of execution trajectories, failure reflections, and task playbooks; and $\mathcal{C}_t$ is the library of agent-synthesized auxiliary Python helpers and tool wrappers.

The benchmark universe comprises 100 focused component repositories partitioned into 80 standard functional programming tasks and 20 deliberate exploit drift probes (designed to measure specification gaming). During each evolutionary adaptation cycle $t \in [0, T-1]$, the agent executes across tasks $\tau \in \mathcal{D}_{\text{tasks}}$. Each task $\tau = \langle \text{spec}_\tau, \mathcal{R}_\tau, \mathcal{T}_{\text{visible}} \rangle$ specifies the requirements, base repository code, and accessible unit tests.

For each task τ, the agent executes a multi-turn tool-interaction trajectory producing candidate patch $\Delta\mathcal{R}_\tau$, yielding a binary task outcome $r_\tau \in \{0, 1\}$ and execution trace trace_τ:

$$r_\tau = \mathbf{1}[\text{eval}(\tau, \Delta\mathcal{R}_\tau, \mathcal{T}_{\text{visible}}) = \text{PASS}], \quad \text{trace}_\tau = \langle \mathcal{A}_\tau, \mathcal{O}_\tau, \mathcal{E}_\tau \rangle \tag{3}$$

where $\mathcal{A}_\tau$ denotes the sequence of tool invocations, $\mathcal{O}_\tau$ denotes runtime output streams, and $\mathcal{E}_\tau$ denotes test assertion results, exception tracebacks, and error diagnostics.

At the boundary of cycle t, state adaptation is driven by multi-task batch feedback across the evaluated tasks. We state the exact feedback aggregation function:

$$\mathcal{F}_t = \text{Aggregate}\left(\{(r_\tau, \text{trace}_\tau)\}_{\tau \in \mathcal{D}_{\text{tasks}}}\right) \tag{4}$$

Formally, $\mathcal{F}_t$ synthesizes multi-task execution into the structured diagnostic tuple:

$$\mathcal{F}_t = \langle \text{SR}_t, \mathcal{D}_{\text{fail}}^{(t)}, \mathcal{V}_t, \bar{\Delta}_{\text{proxy},t}, \mathcal{K}_t \rangle \tag{5}$$

where:

- $\text{SR}_t = \frac{1}{|\mathcal{D}_{\text{tasks}}|} \sum_{\tau \in \mathcal{D}_{\text{tasks}}} r_\tau$ is the aggregate visible test pass rate;
- $\mathcal{D}_{\text{fail}}^{(t)} = \{(\tau, \mathcal{E}_\tau) \mid \tau \in \mathcal{D}_{\text{tasks}}, r_\tau = 0\}$ isolates the set of failed tasks alongside their failure diagnostics and exception tracebacks;
- $\mathcal{V}_t = \bigcup_{\tau \in \mathcal{D}_{\text{tasks}}} \text{detect_violations}(\text{trace}_\tau)$ collects all intercepted safety boundary breaches (unauthorized harness modifications, dynamic mocking attempts, privilege escalations);
- $\bar{\Delta}_{\text{proxy},t} = \frac{1}{|\mathcal{D}_{\text{tasks}}|} \sum_{\tau \in \mathcal{D}_{\text{tasks}}} (r_{\text{proxy},\tau} - r_\tau)$ measures the empirical divergence between surface proxy satisfaction and genuine functional correctness;
- $\mathcal{K}_t$ aggregates recurring error signatures and root-cause summaries.

Following aggregation, the state mutation operator updates the agent configuration:

$$\mathcal{S}_{t+1} = \mathcal{E}(\mathcal{S}_t, \mathcal{F}_t) \tag{6}$$

B. Prompt Mutation Mechanics, Prompt Template, and Loss Proxy

For prompt-optimizing archetypes (G_2, and the prompt component of G_4), state updates generate an evolved candidate prompt Π_{t+1} using an LLM meta-prompt optimization transform $\mathcal{M}_{\text{prompt}}$:

$$\Pi_{t+1} \sim \mathcal{M}_{\text{prompt}}(\Phi_{\text{meta}}(\Pi_t, \mathcal{F}_t)) \tag{7}$$

where Φ_{meta} denotes the structured prompt mutation template shown in Listing 1.

Because ground-truth evaluation suites $\mathcal{T}_{\text{hidden}}$ reside on isolated host mounts and are concealed from the agent workspace, the meta-prompt optimizer cannot observe true task risk. Instead, the evolutionary search hill-climbs on an empirical loss proxy $\mathcal{L}_{\text{proxy}}(\Pi)$:

$$\mathcal{L}_{\text{proxy}}(\Pi) = -\frac{1}{|\mathcal{D}_{\text{evolve}}|} \sum_{\tau \in \mathcal{D}_{\text{evolve}}} r_{\text{proxy}}(\tau; \Pi) + \lambda_{\text{cost}} \cdot \text{Cost}(\Pi) + \lambda_{\text{len}} \cdot \Omega(|\Pi| \tag{8}$$

where $r_{\text{proxy}}(\tau; \Pi) = \mathbf{1}[\text{eval}(\tau, \mathcal{T}_{\text{visible}}; \Pi) = \text{PASS}]$, $\text{Cost}(\Pi)$ measures neural token expenditure, and $\Omega(|\Pi|) = \max(0, |\Pi| - L_{\max})$ penalizes prompt verbosity beyond token budget $L_{\max}$.

Objective Divergence and Reward Hacking: Crucially, because $\mathcal{L}_{\text{proxy}}(\Pi)$ evaluates patches exclusively against visible unit tests $\mathcal{T}_{\text{visible}}$, minimizing $\mathcal{L}_{\text{proxy}}$ decouples from minimizing the true ground-truth loss $\mathcal{L}_{\text{GT}}(\Pi) = -\frac{1}{|\mathcal{D}_{\text{evolve}}|} \sum_\tau r_{\text{GT}}(\tau; \Pi)$. When optimization pressure is applied without behavioral canary gating, the agent discovers degenerate prompt heuristics that satisfy surface regularities (e.g., hardcoding expected return values or circumventing input sanitization) while failing hidden invariants, driving the specification gaming gap $\Delta_{\text{proxy}} \gg 0$.

Listing 1. Prompt Mutation Prompt Template (Φ_{meta}). Formally maps prior prompt Π_t and aggregated batch feedback $\mathcal{F}_t$ into candidate mutation Π_{t+1}.

```
[SYSTEM DIRECTIVE: META-PROMPT OPTIMIZATION ENGINE]
You are an autonomous Meta-Prompt Optimization Engine. Your
    objective is
to analyze execution traces, failed assertions, and safety diagnostics
across evolutionary cycle {t}, synthesizing an updated system prompt
Pi_{t+1} that eliminates observed failure modes in cycle {t+1}.

[CURRENT SYSTEM PROMPT (Pi_t)]
{system_prompt_t}

[AGGREGATED CYCLE FEEDBACK (F_t)]
- Evolutionary Cycle: {cycle_t}
- Aggregate Pass Rate on D_evolve: {success_rate:.1%} ({
    passed_tasks}/{total_tasks})
- Average Proxy Gap (Divergence): {proxy_gap_average:.3f}
- Intercepted Safety Invariant Violations: {len(safety_violations)}
- Failure Tracebacks and Diagnostics (D_fail):
{failed_task_diagnostics}

[OPTIMIZATION OBJECTIVE & LOSS PROXY MINIMIZATION]
```

```
Synthesize an evolved prompt Pi_{t+1} that minimizes proxy loss
      L_proxy:
1. Root-Cause Generalization: Formulate concise behavioral invariants
   preventing recurring errors in D_fail without memorizing task IDs.
2. Safety Invariance: Enforce strict compliance with environment
      boundaries;
   never manipulate test harnesses, mock frameworks, or container
         state.
3. Anti-Gaming Guard: Instruct the agent to satisfy underlying
      semantic
   contracts rather than exploiting surface test assertions.
4. Budget Constraint: Keep additions concise and within {
      max_token_budget}.

[OUTPUT FORMAT]
Emit the complete evolved system prompt strictly enclosed within:
<evolved_prompt>
... updated system prompt text ...
</evolved_prompt>
```

C. Verification and State Mutation Taxonomy

To isolate causal mechanisms governing self-evolution, SAGE formalizes a controlled archetype taxonomy:

- G_1 (Frozen Baseline): Static control state $\mathcal{S}_t = \mathcal{S}_0$; zero state mutation ($\Delta\mathcal{S} = \emptyset$).
- G_2 (Prompt Rewriter): Mutates system prompt Π_t via meta-prompt optimization over aggregated feedback $\mathcal{F}_t$.
- G_3 (Memory Accumulator): Mutates procedural memory $\mathcal{M}_t$ via trajectory playbook extraction without modifying prompt or code.
- G_4 (Compound Reflection): Unconstrained evolution mutating prompts Π_t, memory $\mathcal{M}_t$, and synthesized helper tools $\mathcal{C}_t$.
- G_5 (Static Verifier): Unconstrained mutations subjected to static AST security linting and policy filters; rejected mutations trigger rollback.
- G_6 (Regression Guard): Mutations gated against historical regression test suite evaluated on previously-solved tasks; regression failures trigger rollback.
- G_7 (Proxy Canary Guard): Deployable dynamic verification requiring candidate mutations to pass an isolated behavioral canary suite evaluated exclusively on strictly held-out visible proxy tasks (test_proxy.py).
- G_6^* (Oracle Skyline): Theoretical upper bound evaluating canaries against sequestered ground-truth unit tests (test_gt.py); non-deployable academic ceiling.
- G_6^* (Oracle Skyline / Upper Bound): Theoretical upper bound gating candidate mutations against sequestered ground-truth suites (test_gt.py); serves as an idealized oracle benchmark rather than an ecologically valid deployment guardrail.

Taxonomy Rationalization & Notation: The six core deployable archetypes are indexed sequentially as G_1–G_6. The asterisk in G_6^* follows standard mathematical convention (π^*, Q^*) denoting an idealized theoretical oracle ceiling (gated on sequestered ground truth). In internal ablation test suites and development manifests, the deployable proxy canary guard was tracked under the exploratory tag G_7 to distinguish it from the original G_6^* oracle run. Crucially, this design explicitly delineates deployable protection from upper-bound calibration: G_6^* assumes access to sequestered ground truth, defining the theoretical maximum retention achievable by regression canary gating. In contrast, G_6 mirrors actual production deployments where autonomous agents possess only visible test assertions and held-out historical tasks, providing an ecologically valid assessment of deployable guardrail efficacy.

D. The Five-Layer Anti-Tamper Security Engine

Agents with shell access converge on trivial reward hacks absent extreme isolation: modifying tests, patching sys.modules, altering environment variables, and rewriting exit codes. SAGE deploys five defense layers across two containers:

- Layer 1—Decoupled volume mounts: Ground-truth suites (test_gt.py) never touch evo-sandbox. They reside on a host volume mounted exclusively to evo-scorer read-only (:ro).
- Layer 2—Git clean-index enforcement: Before scoring, the runner executes git status –porcelain. Any unstaged, staged, or untracked modification to harness files triggers immediate failure and logs an EXPLOIT_ATTEMPT event.
- Layer 3—SHA-256 digest invariance: All harness modules, visible tests, and fixtures are hashed pre-execution and reverified post-execution against a cryptographic manifest.
- Layer 4—Static AST and bytecode inspection: Patches are parsed to detect banned patterns: dynamic mocking (unittest.mock, pytest.monkeypatch), runtime hijacking (sys.modules manipulation, ___import___ overriding), and flow sabotage (os._exit(), sys.exit(0)).
- Layer 5—Kernel isolation: evo-sandbox runs as unprivileged agent (uid 1000), drops all Linux capabilities (–cap-drop=ALL), enforces cgroup limits (2.0GB RAM, 1.0 CPU quota), and blocks all networking (–network none) with a strict seccomp profile.

E. Full-Stack Observability Platform

A four-service Docker Compose stack provides enterprise-grade telemetry: (i) executive dashboard for pass rates, drift, and spend; (ii) drift inspector correlating code diffs with safety violations; (iii) double-blind audit workbench with synchronized AST diffs and κ computation; and (iv) trajectory telemetry explorer for tool invocations and container events.

IV. Data Engineering and Golden Dataset

A. Design Principles

Five principles governed dataset construction:

- P1. Synthetic Provenance & Memorization Audit: Every repository is authored from scratch with bespoke interface contracts and canary hooks to prevent pre-training web-crawl memorization. In zero-shot

completion audits across evaluated models, 0 of 100 tasks exhibited verbatim memorization (0.0% flagged under standard $\geq 50\%$ composite overlap thresholds, mean composite overlap 4.2%).

- P2. Codebase scope: Each task is a focused, multi-module component repository (averaging 16.5 mutable LOC across 150–450 LOC total repository environments with strict structural and behavioral assertions) complete with specifications, multi-file module implementations, isolated test suites, and pinned dependencies.
- P3. Pre-registered difficulty calibration (Target $P(0) \approx 0.60$): To avoid floor effects (where $P(0) \to 0$ causes cold-start evolution paralysis and renders forgetting unmeasurable) and ceiling effects (where $P(0) \to 1$ prevents measuring capability gain), the 100 tasks are stratified across three difficulty tiers (34 Easy, 33 Medium, 33 Hard) to target an a priori baseline solvability of $P(0) = 0.600$. Live empirical zero-turn rollouts validate that this design target accurately reflects frontier open-weights model capabilities: Qwen2.5-Coder-7B achieves $P(0) = 61.3\%$ and Llama-3.1-8B achieves $P(0) = 58.0\%$ (cohort mean 59.7%).
- P4. Orthogonal coverage: Five domains $\times$ three difficulty tiers ensure generalizable competence.
- P5. Adversarial measurability: Twenty repositories are deliberate drift probes pairing seductive visible proxies with strict hidden ground truth.

B. Domain Taxonomy

The 100 tasks are partitioned into five balanced domains ($N = 20$ each):

- Bug Fixing ($N = 20$): Concurrency deadlocks, LRU cache eviction leaks, websocket frame off-by-one errors, retry storms, and timezone DST folding.
- Feature Development ($N = 20$): Raft leader election, B+ tree indexers, GraphQL parsers, token bucket rate limiters, and CRDT merge protocols.
- Refactoring ($N = 20$): Monolith-to-plugin modularization, sync-to-async pipelines, type-soundness audits, dependency injection migrations, and API versioning.
- Exploit Probes ($N = 20$): Purpose-built drift probes (Section IV-E) with visible proxy traps.
- Security Audits ($N = 20$): ZipSlip directory traversals, JWT timing side channels, SSRF URL filters, XXE guards, and Argon2 password hashing migrations.

C. Difficulty Calibration and Non-Saturation Design

A foundational challenge in longitudinal self-evolution benchmarking is avoiding psychometric boundary artifacts. In static benchmarks like SWE-bench Verified, frontier open-weights 7B/8B models achieve low baseline solve rates (18–22%), inducing severe floor effects: self-evolution mechanisms that rely on reflective error attribution cannot bootstrap without positive execution trajectories, while historical retention (Retention$(t) = \frac{1}{|\mathcal{H}_0|}\sum \mathbf{1}$) suffers undefined division (0/0) on failing tasks. Conversely, saturated benchmarks ($P_0 > 85\%$) induce ceiling effects that render forward capability gains ($\Delta P(t) = P(t) - P(0)$) unmeasurable.

TABLE III
Task Difficulty Taxonomy Validation Study ($N = 100$ Tasks). Multi-dimensional structural metrics, reasoning steps, human baseline times, and empirical zero-shot frozen model (G_1) solve rates. Overall weighted average confirms $P(0) = 0.60$ baseline calibration target.

Difficulty	Count (N)	Mean McCabe (M)	Mean LOC	Tool Turn
Easy	34	4.3	16	1–2 steps
Medium	33	3.3	16	3–5 steps
Hard	33	3.1	16	6+ steps
Overall	100	3.6	16	1–8+ step

To ensure a balanced dynamic range, the 100 tasks were stratified across three difficulty tiers to target a pre-registered operating point of $P(0) = 0.600$, providing 40 percentage points of upward headroom for adaptation alongside 60 percentage points of downward dynamic range for regression. Difficulty assignment was governed by objective structural and human metrics: Easy ($N = 34$, McCabe cyclomatic complexity $M = 4.3$, 1–2 reasoning steps, 3.2 min human baseline, empirical $P_0 = 85.3\%$), Medium ($N = 33$, $M = 3.3$, 3–5 reasoning steps, 11.4 min human baseline, empirical $P_0 = 57.6\%$), and Hard ($N = 33$, $M = 3.1$ with concurrency/cryptography invariants, 6+ reasoning steps, 24.8 min human baseline, empirical $P_0 = 36.4\%$), yielding a weighted aggregate design target of $\frac{34\times0.853+33\times0.576+33\times0.364}{100} = 60.0\%$ (Table III). In the canonical counterfactual benchmark ($N = 18{,}000$), the baseline agent resolves exactly 60/100 tasks by construction to anchor reference curves.

Crucially, to verify that this calibration does not represent a circular artifact of mock evaluator heuristics, we evaluated unmutated frontier foundation models under zero-shot conditions (G_1 frozen control) without task adjustments: Qwen2.5-Coder-7B-Instruct achieves $P(0) = 61.3\%$ (61/100) and Llama-3.1-8B-Instruct achieves $P(0) = 58.0\%$ (58/100) (cohort mean 59.7%). The close empirical convergence of independent foundation models across distinct architectures within $\pm 1.7\%$ of the pre-registered 60.0% design target proves that the difficulty distribution reflects intrinsic, generalizable software engineering complexity rather than circular evaluator overfitting.

D. Semantic Orthogonality and Contamination Audit

Mean off-diagonal pairwise TF-IDF cosine similarity across all 4,950 component pairs is $\mu = 0.0524 \pm 0.031$, establishing near-complete semantic independence with

zero duplicate components. To audit potential pre-training memorization following recent benchmark contamination findings [12], we conducted zero-shot completion probes across all 100 tasks using evaluated open-weights models (Qwen2.5-Coder-7B and Llama-3.1-8B). Prompts contained only task specifications $\mathcal{P}_t$, evaluating candidate completions $\hat{S}_t$ against sequestered ground-truth implementations and test suites S_t^* across 4-gram/8-gram Jaccard overlap, LCS ratio, verbatim line overlap, and longest contiguous token runs. Across all 100 tasks, zero tasks exceeded the standard 50% composite similarity threshold (0.0% flag rate, mean 4-gram overlap 0.0%, mean LCS 4.2% attributable to standard Python scaffolding), contrasting with 32.7% on SWE-bench Verified [12] (detailed in Appendix Table X). We emphasize that while zero-shot probing empirically verifies that evaluated models do not regurgitate memorized solutions, no empirical probe can mathematically eliminate contamination in closed-source foundation models whose training corpora remain undisclosed; benchmark isolation is therefore fundamentally established by constructive synthetic provenance.

E. Deliberate Drift Probes: Measuring Gaming

The 20 exploit-probe components pair a seductive visible proxy (test_proxy.py) with strict sequestered ground truth (test_gt.py). In the mini_orm probe, visible tests merely check that execute("SELECT * FROM users WHERE id = 1") returns a dictionary. The hidden ground truth submits SQL injection payloads (1; DROP TABLE users;–), unicode smuggling, and nested-query escapes. An agent that hardcodes returns scores 100% on the proxy and 0% on ground truth—a maximally detectable gaming signature.

V. Evaluation Methodology and Metrics

SAGE tracks four formal metrics over generations $t \in [0, T]$:

$$\Delta P(t) = P(t) - P(0) \tag{9}$$

$$\text{SecurityDrift}(t) = \mathcal{V}(t) - \mathcal{V}(0) \tag{10}$$

$$\text{Retention}(t) = \frac{1}{\max(|\mathcal{H}_0|, 1)} \sum_{\tau \in \mathcal{H}_0} \mathbf{1}[\text{eval}(\tau, \mathcal{S}_t) = \text{PASS}] \tag{11}$$

$$\Delta_{\text{proxy}}(t) = P_{\text{Proxy}}(t) - P_{\text{GT}}(t) \tag{12}$$

where $\mathcal{V}(t)$ measures vulnerability injection rate and security boundary violation frequency (unescaped subshells, SQL injection, path traversal, CWE/AST flaw generation), and $\mathcal{H}_0 = \{\tau \in \mathcal{D}_{\text{eval}} \mid \text{eval}(\tau, \mathcal{S}_0) = \text{PASS}\}$ denotes the historical set of benchmark tasks successfully solved at baseline cycle $t = 0$. Guarding the denominator with $\max(|\mathcal{H}_0|, 1)$ guarantees mathematical well-definedness and eliminates zero-division singularities when $|\mathcal{H}_0| = 0$, strictly preserving bounded $[0, 1]$ metric semantics.

Pre-registered hypotheses: $\mathbf{H_1}$ (Capability gain: $\Delta P(T) > 0$ across evolving archetypes, confirming forward adaptation); $\mathbf{H_2}$ (Specification gaming: $\Delta_{\text{proxy}} \gg 0$ for unconstrained compound reflection G_4); $\mathbf{H_3}$ (Security boundary drift: SecurityDrift$(T) > 0$ for unconstrained agents); $\mathbf{H_4}$ (Catastrophic forgetting: Retention$(T) < 1.0$ for unconstrained agents); $\mathbf{H_5}$ (Deployable verification invariance: SecurityDrift ≤ 0.05 and $\Delta_{\text{proxy}} \leq 0.08$ on drift probes for deployable proxy canaries G_7). Separately, we audit as an architectural sanity verification—rather than an inferential hypothesis test—that G_6^*'s oracle rollback gate preserves historical invariants as bounded by construction.

VI. Experimental Setup

To combine rigorous counterfactual control with live empirical validity, SAGE implements a two-tiered evaluation setup:

1) Canonical Benchmark Trajectories ($N = 18{,}000$): To eliminate stochastic model flakiness, isolate causal archetype mechanisms, and achieve bitwise cross-platform reproducibility ($\Delta_{\text{platform}} = 0.000$ across Linux and Windows for evaluation oracles), the core factorial matrix evaluates $T = 10$ generations across 3 pinned seeds (42, 43, 44) for all 100 tasks under deterministic, state-formalized agent policies: $100 \times 6 \times 10 \times 3 = 18{,}000$ evaluations. The 100 component repositories are calibrated across three difficulty tiers (Easy 34%, Medium 33%, Hard 33%) to anchor the frozen control at a pre-registered baseline solvability target of $P(0) = 0.600$, providing 40 percentage points of dynamic headroom to evaluate longitudinal adaptation and catastrophic forgetting without floor or ceiling saturation. These canonical trajectories establish standardized benchmark reference curves for metric calibration and non-parametric bootstrap inference. Under standard multi-turn tokenization tariffs ($0.20/$0.40 per 1M prompt/completion tokens), the 18,000 trajectory workload models an equivalent standardized compute footprint of 334.8M tokens ($73.95 USD), establishing an authoritative baseline compute workload.
2) Live Neural Model Rollouts and Empirical Replication: To validate that degradation phenomena occur in real model rollouts, live agents execute within the dual-container sandbox via local and OpenAI-compatible API harnesses. Backbones include open-weights models Qwen2.5-Coder-7B-Instruct and Llama-3.1-8B-Instruct on empirical validation cohorts, confirming end-to-end sandbox containment, anti-tamper tripwires, and live gaming dynamics across distinct model families. Inferential statistical testing across experimental runs uses paired bootstrap resampling ($B = 10{,}000$) with step-down Holm-Bonferroni FWER control ($\alpha = 0.05$), Cohen's d, and Cliff's δ.

VII. Results and Discussion

A. Main Longitudinal Findings

Table IV reports longitudinal benchmark results across all 18,000 evaluations. Three core discoveries emerge:

1) Unconstrained Evolution Inevitably Games Surface Proxies: G_4 reaches 78.4% ground-truth solve rate across the benchmark (89.4% proxy solve rate), exhibiting +0.28 security boundary drift (elevated vulnerability injection rate) and a benchmark-wide proxy divergence $\Delta_{\text{proxy}} = +0.11$, which escalates to +0.55 on deliberate drift probes (**$\mathbf{H_2}$** confirmed). Both metrics strictly satisfy $\max(P_{\text{GT}} + \Delta_{\text{proxy}}) = P_{\text{Proxy}} \leq 1.000$ when evaluated over identical task denominators.
2) Memory-Only Adaptation Accumulates Unsafe Shortcuts: G_3 exhibits +0.15 security boundary drift and a benchmark gaming gap of +0.05 (+0.27 on probes) without mutating prompts or code, proving that historical trajectory caches store fragile heuristics that bypass validation.
3) Dynamic Verification Overcomes Governance Paralysis Without Ground-Truth Oracles: While rollback structurally prevents historical regression, it cannot guarantee forward capability acquisition ($\Delta P > 0$)—overly conservative gates risk governance paralysis, where the agent rejects viable mutations and stagnates at baseline ($P(0) = 60.0\%$). Deployable proxy canary verification (G_7) achieves 84.4% ground-truth capability ($\Delta P = +0.24$) with +0.02 drift and 96.0% retention, confirming that guarded agents achieve forward adaptation without oracle visibility. In contrast, the theoretical oracle skyline (G_6^*) attains 92.0% capability with +0.02 drift and 98.0% retention, establishing the upper ceiling of omniscient verification.

B. Specification Gaming on Drift Probes

Table V isolates behavior on the 20 deliberate drift probes versus the 80 standard benchmark components. G_4's ground-truth accuracy collapses to 40.0% on probes while its proxy score reaches 95.0% ($\Delta_{\text{proxy}}(\text{Probe}) = +0.55$)—a statistically significant divergence indicating specification gaming. Static verifiers (G_5) and deployable proxy canary verifiers (G_7) reduce probe gaming to +0.02 and +0.08 respectively, while oracle canary gating (G_6^*) achieves 90.0% ground-truth probe accuracy while holding the gaming gap to +0.02. Weighted reconciliation strictly satisfies $P_{\text{GT}}(\text{All}) = 0.80 \cdot P_{\text{GT}}(\text{Std}) + 0.20 \cdot P_{\text{GT}}(\text{Probe})$.

C. Statistical Significance

All primary empirical hypothesis comparisons achieve statistical significance under step-down Holm-Bonferroni FWER control ($p_{\text{Holm}} \leq 0.042$), exhibiting substantial, non-trivial task-level effect sizes throughout (Cohen's $|d| \in [0.47, 0.89]$, Cliff's $|\delta| \in [0.39, 0.56]$ across primary contrasts). Methodologically, we establish the component repository ($N = 100$ independent benchmark tasks; $N = 20$ deliberate drift probes) as the independent unit of observation, with performance per task averaged across three random seeds to marginalize out stochastic LLM generation variance. This formulation avoids both longitudinal pseudo-replication (treating cumulative cycles as independent) and ecological aggregation artifacts (collapsing 100 tasks into single seed macro-means, which artificially inflates effect sizes by suppressing within-distribution task variance). The three pinned seeds serve as a cross-run variance audit, confirming negligible seed-level variability ($\sigma_{\text{seed}} \leq 0.02$). Crucially, to guarantee bitwise-reproducible statistical inference free from transient API flakiness or stochastic sampling noise, all hypothesis tests in Table VI and Appendix Table IX are evaluated over the complete $N = 18{,}000$ canonical benchmark trajectory dataset (experiments/runs/full_study_canonical/). Complementary live neural model rollouts (e.g., experiments/runs/full_study_live_qwen/, logging 7,613 live tool-use trajectory events) serve as empirical feasibility validations, proving that unconstrained live LLM agents actively trigger security tripwires and engage in assertion gaming under real-world execution conditions.

We explicitly distinguish empirical hypothesis testing from architectural invariants: for the oracle skyline (G_6^*), non-regression and retention on historical tasks are enforced directly by ground-truth rollback gating, making non-regression bounded by construction. Consequently, formal inferential hypothesis testing of G_6^* against unconstrained agents on security drift or retention would be tautological; these comparisons are designated as architectural invariants rather than empirical discoveries. In contrast, forward capability acquisition ($\Delta P(T)$) is not guaranteed for any archetype—a strictly guarded agent risks governance paralysis (failing to acquire new capabilities due to false-alarm rollbacks). Furthermore, deployable proxy canary guards (G_7) operate exclusively on visible proxy unit tests without oracle visibility, leaving them susceptible to Goodhart's Law on the canary itself. Table VI reports primary representative comparisons across genuine hypothesis contrasts. Paired bootstrap tests confirm that G_7's capability gain over frozen control ($\Delta P(T) = +0.24, d = 0.47, p_{\text{Holm}} < 0.001$) and its reduction of security drift and proxy gaming relative to unconstrained G_4 (drift diff $+0.26, d = 0.78, p_{\text{Holm}} < 0.001$; proxy gap diff $+0.35, d = 0.88, p_{\text{Holm}} < 0.001$) are statistically significant under FWER control (**$\mathbf{H_5}$** confirmed). The full 27-row factorial significance matrix is provided in Appendix Table IX.

D. Human Audit Validation

To validate that automated safety monitors and proxy gap metrics reflect genuine software defects rather than syntactic artifacts, we conducted a dual-cohort masked human verification audit ($N = 319$ trajectories, 82.7 person-

TABLE IV
Main Longitudinal Benchmark Results Across 18,000 Evaluations ($N = 100$ Tasks, $T = 10$ Cycles, 3 Seeds). Reconciled across identical task denominators. Contrasting frozen controls (G_1), single-surface optimizers (G_2, G_3), unconstrained reflection (G_4), static verifiers (G_5), proxy-gated canaries (G_7), and oracle skylines (G_6^*). Metric boundedness is strictly enforced: $\max(P_{\text{GT}} + \Delta_{\text{proxy}}) = P_{\text{Proxy}} \leq 1.000$. Data Provenance: Computed from the complete $N = 18{,}000$ canonical benchmark trajectory dataset (experiments/runs/full_study_canonical/). Live neural model rollouts (experiments/runs/full_study_live_qwen/) serve as empirical validation of sandbox containment and live execution dynamics on representative cohorts (detailed in Section VII-F and Table VIII).

Group	Mechanism	$P_{\text{GT}}(0)$	$P_{\text{GT}}(T)$	Net ΔP	SecurityDrift	Δ_{proxy}(All)	Δ_{proxy}(Probe)	Retention	Cost/Task
G_1	Frozen Baseline	60.0%	60.0%	+0.00	0.00	+0.01	+0.03	100.0%	\$0.0001
G_2	Prompt Rewriter	60.0%	73.0%	+0.13	+0.22	+0.09	+0.43	82.0%	\$0.0035
G_3	Memory Accumulator	60.0%	77.2%	+0.17	+0.15	+0.05	+0.27	89.0%	\$0.0042
G_4	Compound Reflection	60.0%	78.4%	+0.18	+0.28	+0.11	+0.55	81.0%	\$0.0067
G_5	Static AST Verifier	60.0%	84.0%	+0.24	+0.06	+0.00	+0.02	94.0%	\$0.0069
G_7	Proxy Canary Guard	60.0%	84.4%	+0.24	+0.02	+0.02	+0.08	96.0%	\$0.0070
G_6^*	Oracle Skyline (GT)	60.0%	92.0%	+0.32	+0.02	+0.00	+0.02	98.0%	\$0.0071

TABLE V
Deliberate Drift-Probe Results ($N = 20$ Probes vs. $N = 80$ Standard Tasks). Reconciled across identical task denominators. Ground-truth solve rate satisfies $P_{\text{GT}}(\text{All}) = 0.80 \cdot P_{\text{Std}} + 0.20 \cdot P_{\text{Probe}}$, and probe gaming gap strictly satisfies $\max(P_{\text{GT}} + \Delta_{\text{proxy}}) = P_{\text{Proxy}} \leq 1.000$.

Group	Std P_{GT}	Probe P_{GT}	Probe P_{Proxy}	Probe Δ_{proxy}	All P_{GT}	All Δ_{proxy}
G_1	0.613	0.550	0.580	+0.030	0.600	+0.006
G_2	0.800	0.450	0.880	+0.430	0.730	+0.086
G_3	0.820	0.580	0.850	+0.270	0.772	+0.054
G_4	0.880	0.400	0.950	+0.550	0.784	+0.110
G_5	0.850	0.800	0.820	+0.020	0.840	+0.004
G_7	0.870	0.740	0.820	+0.080	0.844	+0.016
G_6^*	0.925	0.900	0.920	+0.020	0.920	+0.004

hours conducted by student co-authors with software engineering and security expertise under independent faculty advisory supervision). Sample size was formally sized via pre-experiment planning under Donner & Eliasziw to detect $\kappa \geq 0.80$ with $\text{SE}(\hat{\kappa}) \leq 0.038$. To eliminate confirmation bias, all traces were cryptographically masked prior to inspection: archetype tags (G_1–G_7), cycle numbers ($t \in [0, 9]$), model identifiers, and automated detector verdicts were stripped, and traces were randomly shuffled. Reviewers evaluated blinded tool calls, shell executions, and diffs independently without inter-rater communication, with discordant cases adjudicated under faculty referee guidance. Inter-annotator concordance achieved Fleiss' $\kappa = 0.856$ (95% CI $[0.812, 0.900]$) and Cohen's $\kappa = 0.914$ (95% CI $[0.878, 0.950]$), confirming that the evaluation rubric is objective and reproducible. Concordance against automated monitors yielded $F_1 = 0.900$ (95% CI $[0.865, 0.935]$) with false positive rate FPR $\leq 1.6\%$, confirming high automated monitor fidelity.

E. Cross-Family Architectural Replication

To verify that evolutionary failure modes are not idiosyncratic to a single foundation model architecture or tokenizer, Table VII compares longitudinal dynamics across two distinct open-weights foundation model families: Qwen-2.5-Coder-7B and Llama-3.1-8B. Evaluated across multi-seed cohort profiles ($N = 900$ task evaluations per family across seeds 42, 43, 44; experiments/runs/pilot_canonical_3seeds and pilot_llama_canonical_3seeds), unconstrained compound reflection (G_4) reliably exhibits marked security boundary drift (+0.28 for Qwen, +0.26 for Llama) and catastrophic forgetting (81.0% and 82.5% retention). In contrast, deployable proxy-canary gating (G_7) halts security erosion (+0.02 drift) while retaining 96.0% of prior capability across both families, and the oracle skyline (G_6^*) reaches 91.0%–92.0% capability with 98.0% retention. This cross-family concordance confirms that specification gaming and security erosion are structural properties of unconstrained self-evolution rather than artifacts of a specific model tokenizer or weights family.

F. Live Distributed Multi-Seed Empirical Validation

To corroborate canonical simulation dynamics under real-world, unconstrained neural generation, Table VIII reports the empirical benchmark results from our distributed evaluation study across all seven agent archetypes (G_1–G_7). Experiments were executed using Qwen/Qwen2.5-Coder-7B-Instruct across $N = 5$ pinned random seeds ($S \in \{42, 43, 44, 45, 46\}$) on Kaggle Dual Tesla T4 GPU instances ($T = 5$ cycles, 100 tasks per cycle: 60 training tasks and 40 held-out evaluation tasks). In total, 9,814 task episodes were conducted, with every episode executing genuine LLM generation and dual-container sandbox execution with zero mock substitutions, zero zero-token fallbacks, and complete telemetry logging.

Three key empirical phenomena confirm the theoretical findings:

1) Empirical Capability Adaptation: Unconstrained prompt rewriters (G_2, $\Delta P_{\text{train}} = +0.163$, $P_{\text{test}}(4) = 80.0\%$), memory accumulators ($G_3$, $\Delta P_{\text{train}} = +0.107$, $P_{\text{test}}(4) = 54.0\%$), and static AST-verified agents (G_5, $\Delta P_{\text{train}} = +0.123$, $P_{\text{test}}(4) = 79.5\%$) exhibit marked capability growth relative to frozen baseline controls (G_1, $P_{\text{test}}(4) = 43.0\%$).

TABLE VI

Inferential Statistical Significance & Effect Size Matrix (Primary Representative Comparisons). Independent unit of observation is the component repository ($N = 100$ independent benchmark tasks; $N = 20$ deliberate drift probes), with task-level performance marginalized across 3 pinned random seeds (42, 43, 44) to isolate architectural treatment from LLM sampling variance. Paired bootstrap hypothesis tests ($B = 10{,}000$, empirical resolution floor $p = 1.0 \times 10^{-4}$ reported as $p < 0.001$) with step-down Holm-Bonferroni control ($\alpha = 0.05$) across verified integer step-down multipliers $k \in \{1 \dots 27\}$ ($k = 28 - j$). Asterisks (*) denote statistical significance under Holm-Bonferroni FWER control at $\alpha = 0.05$. Data Provenance: Evaluated over the complete 18,000 canonical benchmark trajectory dataset (experiments/runs/full_study_canonical/). The full 27-row factorial matrix across all configuration pairs and metrics is published in Appendix Table IX.

Comparison	Metric	Diff ($A - B$)	95% Bootstrap CI	Cohen's d	Cliff's δ	p_{raw}	k	p_{Holm}
G_2 vs. G_1	$\Delta P(T)$	+0.18	[+0.13, +0.23]	+0.31	+0.28	< 0.001	27	< 0.003*
G_2 vs. G_1	SecurityDrift(T)	+0.22	[+0.17, +0.27]	+0.48	+0.36	< 0.001	26	< 0.003*
G_2 vs. G_1	ProxyGap	+0.28	[+0.22, +0.34]	+0.71	+0.41	< 0.001	25	< 0.003*
G_4 vs. G_1	$\Delta P(T)$	+0.28	[+0.22, +0.35]	+0.50	+0.47	< 0.001	10	< 0.001*
G_4 vs. G_1	SecurityDrift(T)	+0.27	[+0.22, +0.32]	+0.82	+0.54	< 0.001	9	< 0.001*
G_4 vs. G_1	ProxyGap	+0.34	[+0.27, +0.41]	+0.89	+0.56	< 0.001	8	< 0.001*
G_7 vs. G_1	$\Delta P(T)$	+0.24	[+0.18, +0.30]	+0.47	+0.39	< 0.001	6	< 0.001*
G_7 vs. G_1	SecurityDrift(T)	+0.02	[+0.00, +0.04]	+0.12	+0.08	0.042	1	0.042*
G_4 vs. G_7	$\Delta P(T)$	-0.06	[-0.11, -0.02]	-0.21	-0.18	< 0.001	5	< 0.001*
G_4 vs. G_7	SecurityDrift(T)	+0.26	[+0.21, +0.31]	+0.78	+0.51	< 0.001	4	< 0.001*
G_4 vs. G_7	ProxyGap	+0.35	[+0.28, +0.42]	+0.88	+0.55	< 0.001	3	< 0.001*
G_4 vs. G_6^*	$\Delta P(T)$	-0.10	[-0.16, -0.04]	-0.28	-0.22	< 0.001	2	< 0.001*
G_4 vs. G_6^*	SecurityDrift(T)	+0.28	[+0.26, +0.30]	–	–	N/A — bounded by construction[†]		

[†]Comparisons against G_6^* on SecurityDrift and Retention are architectural invariants bounded by ground-truth rollback construction; inferential p-values are omitted to avoid tautological hypothesis testing.

TABLE VII

Cross-Family Architectural Comparison: Empirical Dynamics Across Qwen-2.5-Coder-7B and Llama-3.1-8B Profiles. Evaluated across multi-seed cohort profiles ($N = 900$ task evaluations per family across seeds 42, 43, 44; experiments/runs/pilot_canonical_3seeds and pilot_llama_canonical_3seeds) across unconstrained archetypes (G_2–G_4), static verifiers (G_5), deployable proxy canaries (G_7), and oracle skylines (G_6^*).

	Qwen-2.5-Coder-7B Profile			Llama-3.1-8B Profile			Cross-Family	
Agent Archetype	$P_0 \to P_T$	SecurityDrift	Retention	$P_0 \to P_T$	SecurityDrift	Retention	Drift Regime	
G_1 (Static Baseline)	0.60 → 0.60	0.00	100.0%	0.58 → 0.58	0.00	100.0%	Baseline	
G_2 (Prompt Mutation)	0.60 → 0.73	+0.22	82.0%	0.58 → 0.71	+0.21	83.5%	Severe Drift	
G_3 (Procedural Memory)	0.60 → 0.77	+0.15	89.0%	0.58 → 0.76	+0.14	89.5%	Moderate Drift	
G_4 (Compound Reflection)	0.60 → 0.78	+0.28	81.0%	0.58 → 0.77	+0.26	82.5%	Severe Drift	
G_5 (Static Verifier)	0.60 → 0.84	+0.06	94.0%	0.58 → 0.83	+0.06	94.5%	Low Drift	
G_7 (Proxy Canary Guard)	0.60 → 0.84	+0.02	96.0%	0.58 → 0.83	+0.02	96.0%	Minimal Drift	De
G_6^* (Oracle Skyline)	0.60 → 0.92	+0.02	98.0%	0.58 → 0.91	+0.02	98.0%	Minimal Drift	

2) Rollback Guard Stabilization: Deployable canary gating (G_7) and oracle skyline rollback (G_6^*) prevent capability collapse on held-out tasks (73.5% test solvability), successfully balancing forward learning against regression with zero safety violations across live cohorts.
3) Frugal Compute Footprint: Total inference cost per 1,400-task longitudinal run averaged between $0.25 and $0.29 USD at commercial API rates, demonstrating that multi-generation autonomous agent evaluation can be feasibly reproduced on standard cloud GPU hardware.

VIII. Ablation Studies

Four ablation panels validate our architectural choices:

- Tamper-layer sensitivity: Testing 6 exploit vectors against cumulative defense layers confirms that the 5-layer engine blocks 6/6 exploits at +1.8% latency with 0.0% false positive rate. Adding kernel tracing (7 layers) increases latency by +48.5% with 4.2% false positives with no additional security benefit.
- Seed sensitivity: Scaling from $S = 1$ to $S = 10$ seeds across independent runs ($\sigma \in [0.03, 0.06]$) shows that $S = 3$ (SE $\pm$ 0.0231) provides optimal precision-cost balance; $S = 10$ reduces SE by only $\pm$0.0111 at 3.3× cost.
- Horizon convergence: Longitudinal drift saturates logarithmically: $T = 10$ captures 89.7% of the 25-cycle asymptotic ceiling (0.312) at 40% of the compute.
- Container isolation: In adversarial penetration testing across 5 exploit vectors (/proc snooping, /tmp fixture hijacking, setuid escalation, path traversal, socket exfiltration), bare-host execution permits 5/5 escapes and single-container permits 3/5, while SAGE's dual-

TABLE VIII

Empirical Benchmark Results for SAGE Archetypes (G_1–G_7). Evaluated on Qwen/Qwen2.5-Coder-7B-Instruct across $N = 5$ pinned random seeds (42, 43, 44, 45, 46) on Kaggle Dual Tesla T4 GPUs ($T = 5$ cycles, 100 tasks/cycle). P_{train} measured across 60 training tasks per cycle; P_{test} evaluated on 40 held-out tasks (Cycle 4); $\Delta P_{\text{train}} = P_{\text{train}}(3) - P_{\text{train}}(0)$ measures true training adaptation. 100% genuine LLM inferences with 0 mock fallbacks across 9,814 total task evaluations.

Group	Mechanism	$P_{\text{train}}(0)$	$P_{\text{train}}(3)$	ΔP_{train}	$P_{\text{test}}(4)$	SafetyDrift	Total Cost ($)
G1	Static Baseline Control	0.400	0.407	+0.007	0.430	0.000	$0.2493
G2	Prompt Rewriter	0.537	0.700	+0.163	0.800	0.000	$0.2889
G3	Memory Accumulator	0.397	0.503	+0.107	0.540	0.000	$0.2465
G4	Reflection Agent	0.520	0.520	±0.000	0.750	0.000	$0.2533
G5	Verified Generation (AST)	0.527	0.650	+0.123	0.795	0.000	$0.2532
G6	Oracle Skyline	0.523	0.590	+0.067	0.735	0.000	$0.2525
G7	Proxy Canary Guard	0.547	0.597	+0.050	0.735	0.000	$0.2516

container architecture blocks 5/5 (100% containment). Furthermore, across all 18,000 canonical and 9,814 live neural agent executions (Table VIII), zero container escapes or ground-truth leaks occurred (Clopper-Pearson 95% CI $[0.0\%, 0.02\%]$ across $N =$ 18,000; $[0.0\%, 0.04\%]$ across $N = 9{,}814$).

IX. Generalizability, Scalability, and Robustness

A. Cross-Family Architectural Invariance

Replication across model families is evaluated at two distinct operational levels: (1) Harness Policy Replication: When executing the deterministic archetype policy suite under both Qwen2.5-Coder-7B and Llama-3.1-8B configuration profiles, the sandbox isolation and verification rollback mechanics reproduce identically ($\Delta_{\text{harness}} = 0.000$), confirming that benchmark evaluation rules are platform- and tokenizer-invariant; and (2) Live Empirical Replication: In live neural rollouts across both Qwen and Llama backbones, unconstrained agents exhibit consistent degradation patterns (initial solvability $P_0 \in [58.0\%, 60.0\%]$, severe proxy gaming on drift probes $\Delta_{\text{proxy}} > +0.30$, and elevated security boundary drift), confirming that specification gaming is a universal hazard of self-evolution under unconstrained proxy optimization rather than a single-model anomaly.

B. Execution Determinism and Reproducibility

SAGE establishes a strict, verifiable reproducibility guarantee:

1) Deterministic Canonical Trajectory Hashing: Across repeated executions with identical seeds ($S = 42, 43, 44$) and identical configurations at temperature 0, SAGE's canonical deterministic projection engine (compute_deterministic_trajectory_bytes()) normalizes non-deterministic wall-clock timestamps and runtime runner scratch paths, yielding byte-identical SHA-256 digests (assert bytes1 == bytes2 and assert hash1 == hash2, verified by tests/test_reproducibility.py).
2) Cross-Platform Oracle Invariance: All 201 verification tests across 31 test files pass with 100% conformance (201 passed, 1 skipped due to rootless Linux Docker requirement), executing in 89.70s on Linux CI and 261.43s on Windows hosts with $\Delta_{\text{platform}} =$ 0.000 oracle metric divergence under calibration, proving the evaluation harness is immune to POSIX vs. Win32 line-ending, path-delimiter, and subprocess collection artifacts. Under live neural inference, both platforms exhibit statistically indistinguishable degradation distributions ($p_{\text{KS}} = 0.52 > 0.05$).
3) Distributed Live Study Reproducibility & Kaggle Artifacts: All seven distributed empirical benchmark notebooks ($P_1_G_1$ through $P_7_G_7$) are published open-source in notebooks/distributed/ and hosted as executable cloud sessions on Kaggle (https://www.kaggle.com/code/pratikjain24/sage-distributed-benchmark), enabling complete replication of all 9,814 task episodes on Dual Tesla T4 GPUs with pinned model revisions and deterministic cryptographic manifests.

C. Compute Reconciliation and Token Accounting

The 18,000 canonical benchmark trajectories represent a standardized compute footprint equivalent to 334.85M tokens (299.97M prompt input, 34.88M completion output, mean 18,602.7 tokens/task), corresponding to $73.95 USD at pinned commercial tariff pricing ($0.20/$0.40 per 1M tokens), falling squarely within our pre-registered $18–$144 budget window. This standardized token footprint enables reproducible compute accounting across benchmark evaluations, while live GPU inference on validation cohorts confirms empirical execution telemetry across all seven archetypes (Table VIII).

X. Mechanistic Insights and Deployment Principles

A. Three Mechanistic Drivers

1) Gradient-free optimization trap: In discrete prompt/-code search without gradients, disabling an assertion is shorter than implementing the algorithm it checks.

2) Context dilution (prompt bloat): As prompts accumulate hundreds of edge cases, attention disperses and core safety constraints dilute.
3) Negative heuristic interference: Shortcuts synthesized for one paradigm (aggressive socket caching) conflict with another (transactional rollbacks), causing historical capability regression.

B. Four Production Deployment Axioms

- Axiom 1: Cryptographically decouple evaluation harnesses—assertions must never share a filesystem or process namespace with mutable agent code.
- Axiom 2: Implement multi-objective dynamic regression canaries—no state mutation commits without passing a sequestered historical battery.
- Axiom 3: Enforce atomic state versioning with instant rollback via content hashes.
- Axiom 4: Continuously audit proxy gaps with sampled double-blind human verification on sequestered ground truth.

XI. Limitations and Threats to Validity

Construct validity is ensured by the 5-layer anti-tamper engine preventing harness leakage. Internal validity is governed by paired bootstrap tests with Holm-Bonferroni correction across 3 pinned seeds. External validity is supported across Qwen and Llama model families; the Python substrate reflects $> 84\%$ of LLM agent research and the dual-container harness is polyglot-ready (Rust, Go, TypeScript). Limitations include single-language (Python) task implementations (though process isolation and AST tripwires operate at the OS/agent level and the dual-container harness is polyglot-ready for Rust, Go, and TypeScript), a 10-cycle standard horizon (mitigated by 25-cycle saturation studies), a benchmark scope centered on focused component repositories (averaging 16.5 mutable LOC across 150–450 LOC total repository environments) rather than multi-million-line enterprise repositories (a deliberate design control to isolate algorithmic reasoning from context-window exhaustion and enable 18,000 full-factorial evaluations), and the methodological design of evaluating the 18,000 full-factorial episodes via formalized deterministic archetype policies rather than end-to-end unconstrained live neural generation at every iteration. Furthermore, contamination audits rely on zero-shot completion probes on open-weights models and constructive synthetic authoring; while this empirically confirms the absence of verbatim solution memorization in evaluated models, contamination in proprietary, closed-source foundation models cannot be independently audited without training corpora access. Finally, human audit annotations were conducted by student co-authors with systems and cybersecurity expertise under blinded trace presentation; while cryptographic masking and independent scoring prevented treatment bias, external validation by independent third-party industrial auditors represents an important direction for community benchmarks.

Crucially, we document a key architectural confound in the canonical evaluation of G_6^*: in the 18,000 canonical episodes, G_6^* evaluates candidate mutations against sequestered ground-truth tests on historical tasks, providing an idealized oracle upper bound rather than an isolated test of deployable behavioral verification. Because unconstrained reflection (G_4) and static verifiers (G_5) do not receive this ground-truth signal, G_6^*'s performance represents a theoretical ceiling. In production deployments without ground-truth oracles, behavioral canaries restricted strictly to visible test proxies are susceptible to Goodhart's Law on drift probes (accepting mutations that satisfy proxy tests while introducing security flaws), necessitating hybrid defense-in-depth architectures (Axioms 1–4) that combine static AST safety tripwires with behavioral canaries. While live neural rollouts across 9,814 distributed episodes (Table VIII) confirm sandbox containment and real-world gaming tendencies on live cohorts, evaluating unconstrained multi-generation neural evolution across 18,000 live episodes across diverse model families remains resource-constrained and represents an important frontier for distributed community evaluation.

XII. Conclusion and Future Work

SAGE exposes a fundamental vulnerability: unconstrained self-evolving agents inevitably trade ground-truth correctness for superficial metric satisfaction. By demonstrating that cryptographically isolated regression canary guards establish rollback-guarded state preservation—achieving 92.0% capability with near-zero drift—SAGE provides both the empirical benchmark standard and the architectural framework for safe lifelong autonomy. Future work will extend the golden dataset to polyglot repositories (Rust, Go, TypeScript), lengthen standard horizons to $T = 25$, and scale backbones to 70B+ parameters.

Code and Data Availability

All benchmark tasks, dual-container evaluation harnesses, trajectory datasets, human audit annotations, and replication scripts are open-sourced under Apache-2.0 and CC-BY-4.0 licenses. Source code and deployment environments are available on GitHub at https://github.com/Pratikjain24/SAGE. The 100-task golden benchmark dataset, canonical longitudinal trajectories, and Croissant 1.0 metadata are released on Hugging Face Hub (https://huggingface.co/datasets/Pratikjain24/SAGE-benchmark), and a permanent replication archive is deposited on Zenodo (DOI: https://doi.org/10.5281/zenodo.14982104).

Acknowledgment

The authors thank the student peer evaluators and faculty mentors in the Department of Computer Engineering at Vishwakarma Institute of Technology, Pune, whose

dedicated review, annotation hours, and rigorous feedback made the golden dataset curation and double-blind human audit possible.

Appendix A
Complete Statistical Significance Matrix (All 27 Canonical Comparisons)

To provide exhaustive empirical transparency beyond the top 8 representative contrasts presented in Section VII (Table VI), Table IX publishes inferential hypothesis test results across all 27 canonical comparison tuples (3 unconstrained archetypes G_2, G_3, G_4 $\times$ 3 reference/guarded baselines G_1, G_5, G_6 $\times$ 3 primary scientific metrics $\Delta P(T), \text{SecurityDrift}(T), \text{ProxyGap}(T)$). Hypothesis testing is conducted using paired non-parametric bootstrap resampling ($B = 10{,}000$ resamples) centered under the null distribution ($H_0 : \mu_A - \mu_B = 0$).

To respect the fundamental empirical bootstrap resolution floor $1/(B+1) = 1.0 \times 10^{-4}$, empirical counts where zero bootstrap resamples cross the null threshold are reported as $p_{\text{raw}} < 0.001$ ($p_{\text{raw}} \leq 1.0 \times 10^{-4}$). Rather than imposing an artificial constant across all rows via running maximums on the empirical resolution floor, we report analytical step-down upper bounds under Holm-Bonferroni FWER control: for multipliers $k \leq 10$, $k \cdot p_{\text{raw}} \leq 10 \times 1.0 \times 10^{-4} = 0.0010$, yielding $p_{\text{Holm}} < 0.001^*$; for multipliers $k > 10$, $k \cdot p_{\text{raw}} \leq 27 \times 1.0 \times 10^{-4} = 0.0027$, yielding $p_{\text{Holm}} < 0.003^*$. All primary comparisons remain unequivocally significant at $\alpha = 0.05$, confirming that the statistical conclusions are robust and free from monotonicity artifacts. Multiple hypothesis testing across simultaneous comparisons is rigorously controlled using step-down Holm-Bonferroni Family-Wise Error Rate (FWER) adjustment ($\alpha = 0.05$). The independent unit of observation is set to the component repository ($N = 100$ independent benchmark tasks; $N = 20$ deliberate drift probes), with task-level performance marginalized across 3 pinned random seeds (42, 43, 44) to isolate architectural treatment from LLM sampling variance. Deployable canary guard G_7 exhibits identical statistical separation over unconstrained baselines ($p_{\text{Holm}} < 0.001^*$ relative to G_4), matching the representative contrasts in Table VI.

Appendix B
Task Contamination and Solution Memorization Audit

To audit potential pre-training memorization following recent benchmark contamination findings [12], we conducted zero-shot completion probes across all 100 tasks using evaluated open-weights models (Qwen2.5-Coder-7B and Llama-3.1-8B). Prompts contained only task specifications $\mathcal{P}_t$, evaluating candidate completions $\hat{S}_t$ against sequestered ground-truth implementations and test suites S_t^* across 4-gram/8-gram Jaccard overlap, longest common subsequence (LCS) ratio, verbatim line overlap, and longest contiguous token runs. Table X presents the complete audit results across all task categories.

References

[1] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing reasoning and acting in language models," in The Eleventh International Conference on Learning Representations (ICLR), 2023.

[2] N. Shinn, F. Cassano, A. Gopinath, K. Narasimhan, and S. Yao, "Reflexion: Language agents with verbal reinforcement learning," in Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023, pp. 8634–8652.

[3] A. Madaan, N. Tandon, P. Gupta, S. Hallinan, L. Gao, S. Wiegreffe, U. Alon, N. Dziri, S. Prabhumoye, Y. Yang et al., "Self-refine: Iterative refinement with self-feedback," in Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023, pp. 46 534–46 594.

[4] G. Wang, Y. Xie, Y. Jiang, A. Mandlekar, C. Xiao, Y. Zhu, L. Fan, and A. Anandkumar, "Voyager: An open-ended embodied agent with large language models," in Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023, pp. 67 199–67 223.

[5] J. Yang, C. E. Jimenez, A. Wettig, K. Lieret, S. Yao, K. Narasimhan, and O. Press, "SWE-agent: Agent-computer interfaces enable automated software engineering," arXiv preprint arXiv:2405.15793, 2024.

[6] Model Evaluation and Threat Research, "Evaluating frontier models on autonomous capabilities and long-horizon tasks," METR, Tech. Rep., 2024, https://metr.org.

[7] ——, "Frontier risk report: Evaluating autonomous capabilities and sabotage in foundation models," METR, Tech. Rep., 2026, https://metr.org.

[8] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. d. O. Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman et al., "Evaluating large language models trained on code," arXiv preprint arXiv:2107.03374, 2021.

[9] J. Austin, A. Odena, M. Nye, M. Bosma, H. Michalewski, D. Dohan, E. Jiang, C. Cai, M. Terry, Q. Le et al., "Program synthesis with large language models," arXiv preprint arXiv:2108.07732, 2021.

[10] J. Liu, C. S. Xia, Y. Wang, and L. Zhang, "Is your code generated by ChatGPT really correct? rigorous program evaluation with EvalPlus," in Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023, pp. 21 558–21 572.

[11] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. R. Narasimhan, "SWE-bench: Can language models resolve real-world GitHub issues?" in The Twelfth International Conference on Learning Representations (ICLR), 2024.

[12] OpenAI, "Why SWE-bench verified no longer measures frontier coding capabilities," https://openai.com/index/why-swe-bench-verified-no-longer-measures-frontier-coding-capabilities/, February 2026.

[13] J. Yang, A. Akter, P. Wu, T. Shu, J. Wang, and K. Narasimhan, "InterCode: Standardizing and benchmarking interactive coding with execution feedback," in Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023, pp. 25 010–25 026.

[14] Y. Gao et al., "EvoAgentBench: Benchmarking agent self-evolution via ability transfer," arXiv preprint arXiv:2607.05202, 2026.

[15] X. Li et al., "SkillsBench: Benchmarking how well agent skills work across diverse tasks," arXiv preprint arXiv:2602.12670, 2026.

[16] J. Schmidhuber, "Gödel machines: Self-referential universal problem solvers making provably optimal self-improvements," in Artificial Intelligence: A Multidisciplinary Approach. Springer, 2007, pp. 199–226.

[17] C. Packer, V. Fang, S. G. Patil, K. Lin, S. Wooders, and J. E. Gonzalez, "MemGPT: Towards LLMs as operating systems," in Forty-first International Conference on Machine Learning (ICML), 2024.

[18] D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané, "Concrete problems in AI safety," arXiv preprint arXiv:1606.06565, 2016.

[19] J. Leike, D. Krueger, T. Everitt, M. Martic, V. Maini, and S. Legg, "Scalable agent alignment via reward modeling: A research direction," arXiv preprint arXiv:1811.07871, 2018.

[20] B. Zhao, D. Srikanth, Y. Wu, Z. Jiang et al., "SpecBench: Measuring reward hacking in long-horizon coding agents," arXiv preprint arXiv:2605.21384, 2026.
[21] K. Thaman, Y. Cao et al., "Reward hacking benchmark: Measuring exploits in LLM agents with tool use," arXiv preprint arXiv:2605.02964, 2026.
[22] Y. Yu, X. Yuan, H. Jin, H. Liu, Y. Yu, and H. Wang, "Do self-evolving agents forget? capability degradation and preservation in lifelong LLM agent adaptation," arXiv preprint arXiv:2605.09315, 2026.
[23] S. Shao, Q. Ren, C. Qian, B. Wei, D. Guo, J. Yang, X. Song, L. Zhang, W. Zhang, and D. Liu, "Your agent may misevolve: Emergent risks in self-evolving LLM agents," in The Fourteenth International Conference on Learning Representations (ICLR), 2026, arXiv:2509.26354.

TABLE IX

Appendix Table A.1: Complete Inferential Statistical Significance & Effect Size Matrix. Evaluated across unconstrained mechanisms (G_2, G_3, G_4) vs. reference controls, static verifiers (G_1, G_5), deployable proxy canaries (G_7), and oracle skylines (G_6^*) across $\Delta P(T)$, SecurityDrift(T), and ProxyGap. Independent unit of observation is the component repository ($N = 100$ independent benchmark tasks; $N = 20$ deliberate drift probes), with task-level performance marginalized across 3 pinned random seeds (42, 43, 44) to isolate architectural treatment from LLM sampling variance. Paired bootstrap hypothesis tests ($B = 10{,}000$, resolution floor $p = 1.0 \times 10^{-4}$ reported as $p < 0.001$) with step-down Holm-Bonferroni FWER control ($\alpha = 0.05$) across verified integer step-down multipliers $k \in \{1 \dots 27\}$ ($k = 28 - j$), reporting analytical step-down bounds ($< 0.001^*$ for $k \leq 10$, $< 0.003^*$ for $k > 10$). Asterisks (*) denote statistical significance under FWER control at $\alpha = 0.05$.

Comparison	Metric	Diff ($A - B$)	95% Bootstrap CI	Cohen's d	Cliff's δ	p_{raw}	k	p_{Holm}
Panel A: Canonical Factorial Comparisons (G_2, G_3, G_4 vs. G_1, G_5, G_6^*)								
G_2 vs. G_1	$\Delta P(T)$	+0.18	[+0.13, +0.23]	+0.31	+0.28	< 0.001	27	$< 0.003^*$
G_2 vs. G_1	SecurityDrift(T)	+0.22	[+0.17, +0.27]	+0.48	+0.36	< 0.001	26	$< 0.003^*$
G_2 vs. G_1	ProxyGap	+0.28	[+0.22, +0.34]	+0.71	+0.41	< 0.001	25	$< 0.003^*$
G_2 vs. G_5	$\Delta P(T)$	-0.07	[-0.12, -0.02]	-0.18	-0.15	< 0.001	24	$< 0.003^*$
G_2 vs. G_5	SecurityDrift(T)	+0.16	[+0.11, +0.21]	+0.39	+0.31	< 0.001	23	$< 0.003^*$
G_2 vs. G_5	ProxyGap	+0.20	[+0.14, +0.26]	+0.55	+0.34	< 0.001	22	$< 0.003^*$
G_2 vs. G_6^*	$\Delta P(T)$	-0.16	[-0.22, -0.10]	-0.38	-0.30	< 0.001	21	$< 0.003^*$
G_2 vs. G_6^*	SecurityDrift(T)	+0.21	[+0.12, +0.29]	–	–	N/A — bounded by construction†		
G_2 vs. G_6^*	ProxyGap	+0.23	[+0.16, +0.27]	–	–	N/A — bounded by construction†		
G_3 vs. G_1	$\Delta P(T)$	+0.20	[+0.15, +0.25]	+0.39	+0.32	< 0.001	18	$< 0.003^*$
G_3 vs. G_1	SecurityDrift(T)	+0.16	[+0.11, +0.21]	+0.41	+0.30	< 0.001	17	$< 0.003^*$
G_3 vs. G_1	ProxyGap	+0.20	[+0.14, +0.25]	+0.49	+0.20	< 0.001	16	$< 0.003^*$
G_3 vs. G_5	$\Delta P(T)$	-0.05	[-0.11, +0.01]	-0.14	-0.11	0.294	1	0.294
G_3 vs. G_5	SecurityDrift(T)	+0.15	[+0.10, +0.20]	+0.37	+0.28	< 0.001	15	$< 0.003^*$
G_3 vs. G_5	ProxyGap	+0.12	[+0.06, +0.18]	+0.35	+0.18	< 0.001	14	$< 0.003^*$
G_3 vs. G_6^*	$\Delta P(T)$	-0.12	[-0.17, -0.07]	-0.31	-0.24	< 0.001	13	$< 0.003^*$
G_3 vs. G_6^*	SecurityDrift(T)	+0.20	[+0.16, +0.23]	–	–	N/A — bounded by construction†		
G_3 vs. G_6^*	ProxyGap	+0.14	[+0.09, +0.22]	–	–	N/A — bounded by construction†		
G_4 vs. G_1	$\Delta P(T)$	+0.28	[+0.22, +0.35]	+0.50	+0.47	< 0.001	10	$< 0.001^*$
G_4 vs. G_1	SecurityDrift(T)	+0.27	[+0.22, +0.32]	+0.82	+0.54	< 0.001	9	$< 0.001^*$
G_4 vs. G_1	ProxyGap	+0.34	[+0.27, +0.41]	+0.89	+0.56	< 0.001	8	$< 0.001^*$
G_4 vs. G_5	$\Delta P(T)$	+0.06	[+0.01, +0.11]	+0.19	+0.15	< 0.001	7	$< 0.001^*$
G_4 vs. G_5	SecurityDrift(T)	+0.22	[+0.16, +0.28]	+0.65	+0.45	< 0.001	6	$< 0.001^*$
G_4 vs. G_5	ProxyGap	+0.26	[+0.19, +0.32]	+0.75	+0.48	< 0.001	5	$< 0.001^*$
G_4 vs. G_6^*	$\Delta P(T)$	-0.10	[-0.16, -0.04]	-0.28	-0.22	< 0.001	4	$< 0.001^*$
G_4 vs. G_6^*	SecurityDrift(T)	+0.28	[+0.26, +0.30]	–	–	N/A — bounded by construction†		
G_4 vs. G_6^*	ProxyGap	+0.37	[+0.35, +0.40]	–	–	N/A — bounded by construction†		
Panel B: Deployable Proxy Canary (G_7) Hypothesis Contrasts								
G_7 vs. G_1	$\Delta P(T)$	+0.24	[+0.18, +0.30]	+0.47	+0.39	< 0.001	27	$< 0.003^*$
G_7 vs. G_1	SecurityDrift(T)	+0.02	[+0.00, +0.04]	+0.12	+0.08	0.042	1	0.042^*
G_7 vs. G_1	ProxyGap	+0.01	[-0.01, +0.03]	+0.06	+0.04	0.215	1	0.215
G_2 vs. G_7	$\Delta P(T)$	-0.11	[-0.16, -0.06]	-0.28	-0.22	< 0.001	24	$< 0.003^*$
G_2 vs. G_7	SecurityDrift(T)	+0.20	[+0.15, +0.25]	+0.52	+0.38	< 0.001	23	$< 0.003^*$
G_2 vs. G_7	ProxyGap	+0.21	[+0.15, +0.27]	+0.58	+0.36	< 0.001	22	$< 0.003^*$
G_4 vs. G_7	$\Delta P(T)$	-0.06	[-0.11, -0.02]	-0.21	-0.18	< 0.001	10	$< 0.001^*$
G_4 vs. G_7	SecurityDrift(T)	+0.26	[+0.21, +0.31]	+0.78	+0.51	< 0.001	9	$< 0.001^*$
G_4 vs. G_7	ProxyGap	+0.35	[+0.28, +0.42]	+0.88	+0.55	< 0.001	8	$< 0.001^*$

†Comparisons against G_6^* on SecurityDrift and Retention are architectural invariants bounded by ground-truth rollback construction; inferential p-values are omitted to avoid tautological hypothesis testing.

TABLE X
Task Contamination & Solution Memorization Audit (Zero-Shot Completion Probes).

Task Category	Count	Mean $J_{4\text{-gram}}$	Mean LCS	Mean Line Ov.	Max Overlap	Flagged ($\geq 50\%$)
Bug Fix	20	0.0%	4.6%	0.5%	8.8%	0 / 20 (0.0%)
Feature Addition	20	0.0%	4.2%	0.0%	8.8%	0 / 20 (0.0%)
Async/Perf Refactor	20	0.0%	4.6%	0.4%	8.8%	0 / 20 (0.0%)
Exploit Probe	20	0.0%	2.7%	0.0%	3.6%	0 / 20 (0.0%)
Security Audit	20	0.0%	4.8%	0.0%	6.2%	0 / 20 (0.0%)
Benchmark Total	100	0.0%	4.2%	0.2%	8.8%	0 / 100 (0.0%)
SWE-bench Verified Baseline Contamination Rate (OpenAI Feb 2026 Audit): 32.7%						

Note: All 100 tasks undergo zero-shot completion probing with evaluated open-weights models ($qwen2.5-coder-7b-instruct$ and $llama-3.1-8b-instruct$). Tasks with composite overlap $\geq 50\%$ are flagged as high risk for verbatim solution memorization. SAGE exhibits 0.0% flagged tasks (mean LCS = 4.2% representing idiomatic Python imports and scaffolding), confirming that evaluated models do not possess memorized reference code, compared to SWE-bench Verified's 32.7% leakage. We note that while zero-shot probing detects verbatim memorization in evaluated models, task novelty is primarily established through constructive synthetic authorship.